

EM Algorithm

The goal of the EM algorithm is to find maximum likelihood solutions for models having latent variables. Consider a probabilistic model in which we collectively denote all of the observed variables by \mathbf{X} and all of the hidden variables by \mathbf{Z} . The joint distribution $p(\mathbf{X}, \mathbf{Z} \mid \boldsymbol{\theta})$ is governed by a set of parameters denoted $\boldsymbol{\theta}$. Our goal is to maximize the likelihood function that is given by

$$p(\mathbf{X} \mid \boldsymbol{\theta}) = \sum_{\mathbf{Z}} p(\mathbf{X}, \mathbf{Z} \mid \boldsymbol{\theta})$$

This method based on the assumption that the computation of $p(\mathbf{X}, \mathbf{Z} \mid \boldsymbol{\theta})$ is easier than $p(\mathbf{X} \mid \boldsymbol{\theta})$. Let $q(\mathbf{Z})$ be the distribution of \mathbf{Z} . Then is obvious that

$$\ln p(\mathbf{X} \mid \boldsymbol{\theta}) = \sum_{\mathbf{Z}} q(\mathbf{Z}) \ln p(\mathbf{X} \mid \boldsymbol{\theta})$$

We've known that

$$\ln p(\mathbf{X}, \mathbf{Z} \mid \boldsymbol{\theta}) = \ln p(\mathbf{Z} \mid \mathbf{X}, \boldsymbol{\theta}) + \ln p(\mathbf{X} \mid \boldsymbol{\theta})$$

Then

$$\begin{aligned} \ln p(\mathbf{X} \mid \boldsymbol{\theta}) &= \sum_{\mathbf{Z}} \ln q(\mathbf{Z}) p(\mathbf{X} \mid \boldsymbol{\theta}) \\ &= \sum_{\mathbf{Z}} q(\mathbf{Z}) \ln \left\{ \frac{p(\mathbf{X}, \mathbf{Z} \mid \boldsymbol{\theta})}{q(\mathbf{Z})} \right\} - \sum_{\mathbf{Z}} q(\mathbf{Z}) \ln \left\{ \frac{p(\mathbf{Z} \mid \mathbf{X}, \boldsymbol{\theta})}{q(\mathbf{Z})} \right\} \\ &\equiv \mathcal{L}(q, \boldsymbol{\theta}) + \text{KL}(q \parallel p) \end{aligned}$$

Since $\text{KL}(q \parallel p) \geq 0$, $\mathcal{L}(q, \boldsymbol{\theta})$ is the lower bound of $\ln p(\mathbf{X} \mid \boldsymbol{\theta})$, i.e.

$$\mathcal{L}(q, \boldsymbol{\theta}) \leq \ln p(\mathbf{X} \mid \boldsymbol{\theta})$$

Note that $\text{KL}(q \parallel p) = 0$ if and only if $q(\mathbf{Z}) = p(\mathbf{Z} \mid \mathbf{X}, \boldsymbol{\theta})$.

The EM algorithm is a two-stage iterative optimization technique for finding maximum likelihood solutions.

E Step

Let $\boldsymbol{\theta}^{\text{old}}$ be the current parameter vector. Maximize $\mathcal{L}(q, \boldsymbol{\theta})$ with respect to $q(\mathbf{Z})$ with holding $\boldsymbol{\theta}^{\text{old}}$ fixed. $\mathcal{L}(q, \boldsymbol{\theta})$ will reach its maximum $\ln p(\mathbf{X} \mid \boldsymbol{\theta})$ when $q(\mathbf{Z}) = p(\mathbf{Z} \mid \mathbf{X}, \boldsymbol{\theta}^{\text{old}})$. That is,

$$\max \mathcal{L}(q, \boldsymbol{\theta}) = \sum_{\mathbf{Z}} p(\mathbf{Z} \mid \mathbf{X}, \boldsymbol{\theta}^{\text{old}}) \ln p(\mathbf{X}, \mathbf{Z} \mid \boldsymbol{\theta}^{\text{old}}) - \sum_{\mathbf{Z}} p(\mathbf{Z} \mid \mathbf{X}, \boldsymbol{\theta}^{\text{old}}) \ln p(\mathbf{Z} \mid \mathbf{X}, \boldsymbol{\theta}^{\text{old}})$$

M Step

Maximize $\mathcal{L}(q, \boldsymbol{\theta})$ with respect to $\boldsymbol{\theta}$ with holding $q(\mathbf{Z})$ fixed. Here we get new parameter $\boldsymbol{\theta}^{\text{new}}$. This will cause $\mathcal{L}(q, \boldsymbol{\theta})$ increase and then cause $\ln p(\mathbf{X} \mid \boldsymbol{\theta})$ increase subsequently. Furthermore, the $q(\mathbf{Z})$ in this step is determined in the previous step, $\text{KL}(q, \boldsymbol{\theta}^{\text{new}})$ will be nonzero. Hence $\ln p(\mathbf{X} \mid \boldsymbol{\theta})$ increases more than its lowe bound $\mathcal{L}(q, \boldsymbol{\theta})$. From above we have,

$$\begin{aligned} \mathcal{L}(q, \boldsymbol{\theta}) &= \sum_{\mathbf{Z}} p(\mathbf{Z} \mid \mathbf{X}, \boldsymbol{\theta}^{\text{old}}) \ln p(\mathbf{X}, \mathbf{Z} \mid \boldsymbol{\theta}) - \sum_{\mathbf{Z}} p(\mathbf{Z} \mid \mathbf{X}, \boldsymbol{\theta}^{\text{old}}) \ln p(\mathbf{Z} \mid \mathbf{X}, \boldsymbol{\theta}^{\text{old}}) \\ &= \mathcal{Q}(\boldsymbol{\theta}, \boldsymbol{\theta}^{\text{old}}) + \text{constant} \end{aligned}$$

Since $p(\mathbf{Z} \mid \mathbf{X}, \boldsymbol{\theta}^{\text{old}}) = q(\mathbf{Z})$, the last term is a constant. This tells us that we actually maximize the expectation of the complete-data log likelihood.

EM Algorithm

The goal of the EM algorithm is to find maximum likelihood solutions for models having latent variables. Consider a probabilistic model in which we collectively denote all of the observed variables by \mathbf{X} and all of the hidden variables by \mathbf{Z} . The joint distribution $p(\mathbf{X}, \mathbf{Z} \mid \boldsymbol{\theta})$ is governed by a set of parameters denoted $\boldsymbol{\theta}$. Our goal is to maximize the likelihood function that is given by

$$p(\mathbf{X} \mid \boldsymbol{\theta}) = \sum_{\mathbf{Z}} p(\mathbf{X}, \mathbf{Z} \mid \boldsymbol{\theta})$$

This method based on the assumption that the computation of $p(\mathbf{X}, \mathbf{Z} \mid \boldsymbol{\theta})$ is easier than $p(\mathbf{X} \mid \boldsymbol{\theta})$. Let $q(\mathbf{Z})$ be the distribution of \mathbf{Z} . Then is obvious that

$$\ln p(\mathbf{X} \mid \boldsymbol{\theta}) = \sum_{\mathbf{Z}} q(\mathbf{Z}) \ln p(\mathbf{X} \mid \boldsymbol{\theta})$$

We've known that

$$\ln p(\mathbf{X}, \mathbf{Z} \mid \boldsymbol{\theta}) = \ln p(\mathbf{Z} \mid \mathbf{X}, \boldsymbol{\theta}) + \ln p(\mathbf{X} \mid \boldsymbol{\theta})$$

Then

$$\begin{aligned} \ln p(\mathbf{X} \mid \boldsymbol{\theta}) &= \sum_{\mathbf{Z}} \ln q(\mathbf{Z}) p(\mathbf{X} \mid \boldsymbol{\theta}) \\ &= \sum_{\mathbf{Z}} q(\mathbf{Z}) \ln \left\{ \frac{p(\mathbf{X}, \mathbf{Z} \mid \boldsymbol{\theta})}{q(\mathbf{Z})} \right\} - \sum_{\mathbf{Z}} q(\mathbf{Z}) \ln \left\{ \frac{p(\mathbf{Z} \mid \mathbf{X}, \boldsymbol{\theta})}{q(\mathbf{Z})} \right\} \\ &\equiv \mathcal{L}(q, \boldsymbol{\theta}) + \text{KL}(q \parallel p) \end{aligned}$$

Since $\text{KL}(q \parallel p) \geq 0$, $\mathcal{L}(q, \boldsymbol{\theta})$ is the lower bound of $\ln p(\mathbf{X} \mid \boldsymbol{\theta})$, i.e.

$$\mathcal{L}(q, \boldsymbol{\theta}) \leq \ln p(\mathbf{X} \mid \boldsymbol{\theta})$$

Note that $\text{KL}(q \parallel p) = 0$ if and only if $q(\mathbf{Z}) = p(\mathbf{Z} \mid \mathbf{X}, \boldsymbol{\theta})$.

The EM algorithm is a two-stage iterative optimization technique for finding maximum likelihood solutions.

E Step

Let $\boldsymbol{\theta}^{\text{old}}$ be the current parameter vector. Maximize $\mathcal{L}(q, \boldsymbol{\theta})$ with respect to $q(\mathbf{Z})$ with holding $\boldsymbol{\theta}^{\text{old}}$ fixed. $\mathcal{L}(q, \boldsymbol{\theta})$ will reach its maximum $\ln p(\mathbf{X} \mid \boldsymbol{\theta})$ when $q(\mathbf{Z}) = p(\mathbf{Z} \mid \mathbf{X}, \boldsymbol{\theta}^{\text{old}})$. That is,

$$\max \mathcal{L}(q, \boldsymbol{\theta}) = \sum_{\mathbf{Z}} p(\mathbf{Z} \mid \mathbf{X}, \boldsymbol{\theta}^{\text{old}}) \ln p(\mathbf{X}, \mathbf{Z} \mid \boldsymbol{\theta}^{\text{old}}) - \sum_{\mathbf{Z}} p(\mathbf{Z} \mid \mathbf{X}, \boldsymbol{\theta}^{\text{old}}) \ln p(\mathbf{Z} \mid \mathbf{X}, \boldsymbol{\theta}^{\text{old}})$$

M Step

Maximize $\mathcal{L}(q, \boldsymbol{\theta})$ with respect to $\boldsymbol{\theta}$ with holding $q(\mathbf{Z})$ fixed. Here we get new parameter $\boldsymbol{\theta}^{\text{new}}$. This will cause $\mathcal{L}(q, \boldsymbol{\theta})$ increase and then cause $\ln p(\mathbf{X} \mid \boldsymbol{\theta})$ increase subsequently. Furthermore, the $q(\mathbf{Z})$ in this step is determined in the previous step, $\text{KL}(q, \boldsymbol{\theta}^{\text{new}})$ will be nonzero. Hence $\ln p(\mathbf{X} \mid \boldsymbol{\theta})$ increases more than its lowe bound $\mathcal{L}(q, \boldsymbol{\theta})$. From above we have,

$$\begin{aligned} \mathcal{L}(q, \boldsymbol{\theta}) &= \sum_{\mathbf{Z}} p(\mathbf{Z} \mid \mathbf{X}, \boldsymbol{\theta}^{\text{old}}) \ln p(\mathbf{X}, \mathbf{Z} \mid \boldsymbol{\theta}) - \sum_{\mathbf{Z}} p(\mathbf{Z} \mid \mathbf{X}, \boldsymbol{\theta}^{\text{old}}) \ln p(\mathbf{Z} \mid \mathbf{X}, \boldsymbol{\theta}^{\text{old}}) \\ &= \mathcal{Q}(\boldsymbol{\theta}, \boldsymbol{\theta}^{\text{old}}) + \text{constant} \end{aligned}$$

Since $p(\mathbf{Z} \mid \mathbf{X}, \boldsymbol{\theta}^{\text{old}}) = q(\mathbf{Z})$, the last term is a constant. This tells us that we actually maximize the expectation of the complete-data log likelihood.