

Batch Normalization <https://arxiv.org/abs/1502.03167>

Training Time

Let $\{f_1, f_2, f_3, \dots, f_n\}$ be the batched feature vector of some layer (a.k.a. that layer's output, either pre-activation or post-activation). Let's say each $f_i \in \mathbb{R}^d$. The Batch Normalization first computes

$$\begin{aligned}\tilde{f}_i &= (f_i - \mu) / \sigma \\ \mu &= \frac{1}{n} \sum_{i=1}^n f_i \\ \sigma &= \sqrt{\frac{\sum_{i=1}^n (f_i - \mu)^2}{n}}\end{aligned}$$

Second, computes the affine transform

$$\hat{f}_i = \gamma \cdot \tilde{f}_i + \beta$$

where all arithmetic are element-wise, and γ, β are learnable.

We'll use each \hat{f}_i as next layer's input. Each new batch features $\{\hat{f}_i\}_{i=1}^n$ has mean 0 and standard deviation 1.

We can regard BN as another layer takes the whole batch as input, with weight γ and bias β .

Testing Time

During the testing time, mostly we don't have a batched data. Hence no on-line μ, σ to compute. Instead, we use the training data help us get them.

For each i -th batch, we have μ_i, σ_i . We can then use Exponential Moving Average to compute the $\bar{\mu}, \bar{\sigma}$ for testing

$$\begin{aligned}\bar{\mu} &= \alpha \bar{\mu} + (1 - \alpha) \mu_i \\ \bar{\sigma} &= \alpha \bar{\sigma} + (1 - \alpha) \sigma_i\end{aligned}$$

where α is a pre-defined hyper parameter. In pytorch, α is called *momentum* and the default value is 0.1.

Batch ReNormalization

<https://arxiv.org/abs/1702.03275>

Layer Normalization

<https://arxiv.org/abs/1607.06450>

Instance Normalization

<https://arxiv.org/abs/1607.08022>

Group Normalization

<https://arxiv.org/abs/1803.08494>

Weight Normalization

<https://arxiv.org/abs/1602.07868>

Weight Normalization

<https://arxiv.org/abs/1705.10941>